

# Retrospective Approximation (MITACS GRI Internship)

V Pramodh Gopalan, Fabian Bastin

Dept. Informatique et Recherche Opérationnelle  
Université de Montréal

July 28, 2022

# Table of Contents

- 1 Sample Average Approximation, Retrospective Approximation
- 2 L-BFGS
- 3 Common Random Numbers
- 4 Introducing Stopping Tests in Inner Iterations
- 5 Experiments and Results
- 6 Future Work and Improvements

# Sample Average Approximation (SAA)

- Consider the minimization problem:

$$\min_{x \in \mathbb{R}^p} f(x) = \mathbb{E}_{\xi}[F(x, \xi)]$$

- $F$  is deterministic, differentiable, and is relatively easy to compute. It can be non-convex.  $\xi$  is a random vector.  $x$  is of a large dimension.
- Even after knowing  $\xi$ 's distribution, computing  $f(x)$  tough. Hence we approximate function as:

$$\min_{x \in \mathbb{R}^p} \tilde{f}_N(x) = \frac{1}{N} \sum_{i=1}^N F(x, \xi_i)$$

- $\{\xi_i\}_{i=1}^N$  are IID realisations of  $\xi$ .  $\tilde{f}_N(x)$  is now deterministic, can use known techniques. This is Sample Average Approximation (SAA).

- Hope that as  $N$  grows, we converge to actual solution. Special cases include Neural Network training.
- To have theoretical guarantees, need  $N$  to be large. Standard Optimization techniques become infeasible.
- Can we get away with poorer approximations of  $f$  at the beginning?

# Retrospective Approximation(RA)

- RA improves upon SAA by considering a sequence of SAA problems.
- We need a non-decreasing sequence of sample sizes  $\{N_k\}$ , sequence of error-tolerances  $\epsilon_k$ , and a solver capable of solving SAA problems to any desired tolerance level  $\epsilon$ .

Here is the RA algorithm: given an initial guess  $X_0$

- ① At iteration  $k$ , generate an SAA problem with sample size  $N_k$ .
  - ② With  $X_{k-1}$  as initial guess, and  $\epsilon_k$  as the error-tolerance, use the solver to solve the SAA.
  - ③ Obtain the solution  $X_k$ .
- **Disadvantage:** In lower sample sizes, the solver can be prone to overfitting.

- LBFGS (Limited-Memory Broyden–Fletcher–Goldfarb–Shanno) is an optimization algorithm designed to approximate BFGS but using limited memory.
- BFGS approximates the Inverse Hessian, to precondition the gradient to get the descent direction. Cheaper than using true hessian or full matrix approximation.
- BFGS stores dense  $p \times p$  Hessian matrices, which becomes a limitation.
- L-BFGS stores last  $m$  ( $< 10$ ) updates of the iterates and gradients. Lesser memory consumption. These are used to construct inverse Hessian approximations.

# Common Random Numbers (CRN)

- CRN is a simple method for increasing the efficiency of estimating *difference* in performance.
- It introduces dependence in between random variables to reduce variance. Usually, this means the systems be simulated using same stream of random numbers.
- Precisely, given two random variables  $X$  and  $Y$ , and two functions  $f$  and  $g$ , we want to find

$$\mathbb{E}[f(X) - g(Y)]$$

- The effort required to find this is dependant on the variance

$$\mathbf{Var}[f(X) - g(Y)] = \mathbf{Var}[f(X)] + \mathbf{Var}[g(Y)] - 2\mathbf{Cov}[f(X), g(Y)]$$

- if  $X$  and  $Y$  were to be independent,  $\mathbf{Cov}[f(X), g(Y)] = 0$ , but using CRN causes positive correlation, which reduces variance.

# How to construct stopping tests in RA?

- Can we use CRN to find out when a function stops decreasing in value?
- The optimization problem we have is:

$$\min_{x \in \mathbb{R}^p} \tilde{f}_N(x) = \frac{1}{N} \sum_{i=1}^N f(x, \xi_i)$$

- Say we have current iterate  $X_k$  and  $X_{k+1} = X_k + \eta d$ , where  $d$  is a direction of descent yielded by a solver,  $\eta$  being an appropriate step size.



# How to construct Stopping tests in RA?

- In the RA scheme, the subsampled dataset for each iteration is picked at random. So, given a sample size  $N_k$ , we calculate

$$\sigma^2 = \mathbf{Var}[f_{N_k}(X_k) - f_{N_k}(X_{k+1})]$$

- Now, we build an  $\alpha$  confidence interval around  $\mu = \mathbb{E}[f_{N_k}(X_k) - f_{N_k}(X_{k+1})]$  using the variance from earlier step.
- Therefore, if

$$\mu - z_{1-\alpha/2}\sigma > 0$$

we continue the algorithm. Here  $z_\gamma$  returns the  $\gamma$  quantile of the standard Normal distribution.

- These stopping tests take into consideration the noise in gradient, and prevent overfitting in lower sample size problems.

# Comparing RA and SGD

- We use a logistic regression example with 10 covariates and 10,000 samples as a toy dataset to demonstrate that RA can be better than SGD.
- Use the logistic loss function, and set the mini batch sizes and  $\epsilon$  (gradient norm here) by hand.
- Note, we need to use the number of gradient calls as a metric. It serves as a placeholder for amount of work done.
- Even for the best case batch size for SGD, we see that it takes 1.6x more gradient calls than RA (SGD as inner solver).
- One disadvantage is tuning the epsilons by hand, which is time consuming in large scale applications.

# Including Stopping Tests

- In this case, we follow a setup illustrated in [Newton et al., 2021].
- Ordinary Least Squares Regression Problem with 1,000 covariates and 10,000 data samples.
- We use an L-BFGS algorithm, but here the stored history is carried over from one inner iteration to next. We use the Armijo backtracking line search.
- For minibatch scaling, we use a geometric progression, where  $N_{k+1} = 1.2N_k$ , with  $N_0 = 2000$ , for 5 inner iterations.
- In this case, RA performs just as well as SGD, taking approximately the same amount of gradient calls.

# Future Improvements

- Does the sample size have to be non-decreasing? Decreasing sample sizes might help escape saddle points.
- Try stopping tests on non-convex functions.
- Reuse computation to adapt it to online settings.



Newton, D., Bollapragada, R., Pasupathy, R., and Yip, N. K. (2021). Retrospective approximation for smooth stochastic optimization.